


TensorFlow

DEV SUMMIT 2017



TensorBoard



Dandelion Mané
dandelion@google.com



Neural Nets
can be a
black box

TensorBoard is a
flashlight

(but what **is** it)



Write a regex to create a tag group

✕

- ☐ Split on underscores
- ☐ Data download links
- Tooltip sorting method: default



Horizontal Axis

STEP

RELATIVE

WALL

Runs


Write a regex to filter runs

- ☒ lr_1E-05,conv=1,fc=1
- ☒ lr_1E-05,conv=1,fc=2
- ☒ lr_1E-05,conv=2,fc=1
- ☒ lr_1E-05,conv=2,fc=2
- ☒ lr_1E-04,conv=1,fc=1
- TOGGLE ALL RUNS

/usr/local/google/home/dandelion/mnist_tutorial/5


accuracy

accuracy/accuracy



xent

xent/xent_1



Let's see an example



MNIST Architecture



Layer Definition

```
# Define a simple convolutional layer
def conv_layer(input, channels_in, channels_out):
    w = tf.Variable(tf.zeros([5, 5, channels_in, channels_out]))
    b = tf.Variable(tf.zeros([channels_out]))
    conv = tf.nn.conv2d(input, w, strides=[1, 1, 1, 1], padding="SAME")
    act = tf.nn.relu(conv + b)
    return act
```

```
# And a fully connected layer
def fc_layer(input, channels_in, channels_out):
    w = tf.Variable(tf.zeros([channels_in, channels_out]))
    b = tf.Variable(tf.zeros([channels_out]))
    act = tf.nn.relu(tf.matmul(input, w) + b)
    return act
```



Feed-forward Setup

```
# Setup placeholders, and reshape the data
x = tf.placeholder(tf.float32, shape=[None, 784])
y = tf.placeholder(tf.float32, shape=[None, 10])
x_image = tf.reshape(x, [-1, 28, 28, 1])

# Create the network
conv1 = conv_layer(x_image, 1, 32)
pool1 = tf.nn.max_pool(conv1, ksize=[1, 2, 2, 1], strides=[1, 2, 2, 1], padding="SAME")

conv2 = conv_layer(pool1, 32, 64)
pool2 = tf.nn.max_pool(conv2, ksize=[1, 2, 2, 1], strides=[1, 2, 2, 1], padding="SAME")
flattened = tf.reshape(pool2, [-1, 7 * 7 * 64])

fc1 = fc_layer(flattened, 7 * 7 * 64, 1024)
logits = fc_layer(fc1, 1024, 10)
```



Loss & Training

```
# Compute cross entropy as our loss function
cross_entropy = tf.reduce_mean(
    tf.nn.softmax_cross_entropy_with_logits(logits=logits, labels=y))

# Use an AdamOptimizer to train the network
train_step = tf.train.AdamOptimizer(1e-4).minimize(cross_entropy)

# compute the accuracy
correct_prediction = tf.equal(tf.argmax(logits, 1), tf.argmax(y, 1))
accuracy = tf.reduce_mean(tf.cast(correct_prediction, tf.float32))
```



Train the Model

```
# Initialize all the variables
sess.run(tf.global_variables_initializer())

# Train for 2000 steps
for i in range(2000):
    batch = mnist.train.next_batch(100)

    # Occasionally report accuracy
    if i % 500 == 0:
        [train_accuracy] = sess.run([accuracy], feed_dict={x: batch[0], y: batch[1]})
        print("step %d, training accuracy %g" % (i, train_accuracy))

# Run the training step
sess.run(train_step, feed_dict={x: batch[0], y_true: batch[1]})
```



We run our model and...



```
step 0,      training accuracy 10%  
step 500,    training accuracy 12%  
step 1500,   training accuracy 9%  
step 2000,   training accuracy 13%
```

it didn't learn anything!



Let's turn on our flashlight!



Let's visualize the TensorFlow graph.



Visualizing the Graph

```
tf.summary.FileWriter (n):  
    a Python class that writes data for TensorBoard
```

```
writer = tf.summary.FileWriter("/tmp/mnist_demo/1")  
writer.add_graph(sess.graph)
```

```
>> tensorboard --logdir /tmp/mnist_demo/1
```



(DEMO)



Cleaning the Graph

- Node Names
- Name Scopes



Cleaning The Graph

```
def conv_layer(input, channels_in, channels_out, name="conv"):
    with tf.name_scope(name):
        w = tf.Variable(tf.zeros([5, 5, channels_in, channels_out]), name="W")
        b = tf.Variable(tf.zeros([channels_out]), name="B")
        conv = tf.nn.conv2d(input, w, strides=[1, 1, 1, 1], padding="SAME")
        act = tf.nn.relu(conv + b)
        return tf.nn.max_pool(act, ksize=[1, 2, 2, 1], strides=[1, 2, 2, 1], padding="SAME")

def fc_layer(input, channels_in, channels_out, name="fc"):
    with tf.name_scope(name):
        w = tf.Variable(tf.zeros([channels_in, channels_out]), name="W")
        b = tf.Variable(tf.zeros([channels_out]), name="B")
        return tf.nn.relu(tf.matmul(input, w) + b)
```



Cleaning The Graph

```
# Setup placeholders, and reshape the data
x = tf.placeholder(tf.float32, shape=[None, 784], name="x")
x_image = tf.reshape(x, [-1, 28, 28, 1])
y = tf.placeholder(tf.float32, shape=[None, 10], name="labels")

conv1 = conv_layer(x_image, 1, 32, "conv1")
conv2 = conv_layer(conv1, 32, 64, "conv2")

flattened = tf.reshape(conv2, [-1, 7 * 7 * 64])
fc1 = fc_layer(flattened, 7 * 7 * 64, 1024, "fc1")
logits = fc_layer(fc1, 1024, 10, "fc2")
```



Cleaning The Graph

```
with tf.name_scope("xent"):  
    xent = tf.reduce_mean(  
        tf.nn.softmax_cross_entropy_with_logits(logits=logits, labels=y))  
  
with tf.name_scope("train"):  
    train_step = tf.train.AdamOptimizer(1e-4).minimize(xent)  
  
with tf.name_scope("accuracy"):  
    correct_prediction = tf.equal(tf.argmax(logits, 1), tf.argmax(y, 1))  
    accuracy = tf.reduce_mean(tf.cast(correct_prediction, tf.float32))
```



Cleaning the Graph

```
writer = tf.summary.FileWriter("/tmp/mnist_demo/2")  
writer.add_graph(sess.graph)
```

```
>> tensorboard --logdir /tmp/mnist_demo/2
```



(DEMO)



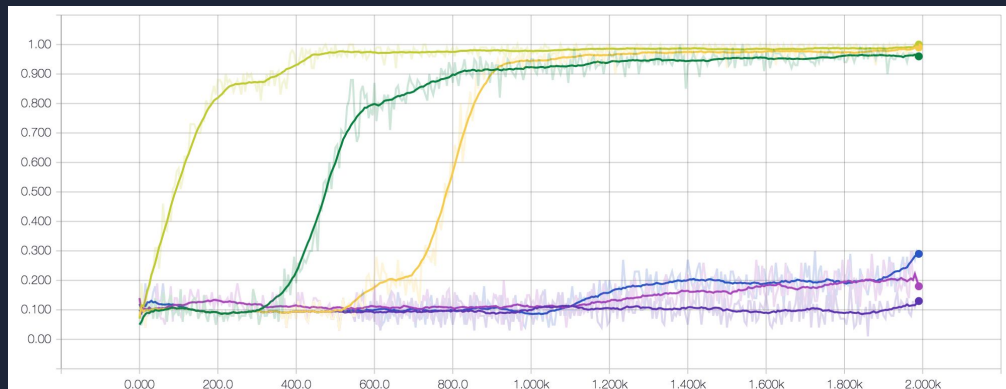
Collect some summaries

`summary(n):`

a TensorFlow op that output protocol buffers containing "summarized" data

Examples:

- `tf.summary.scalar`



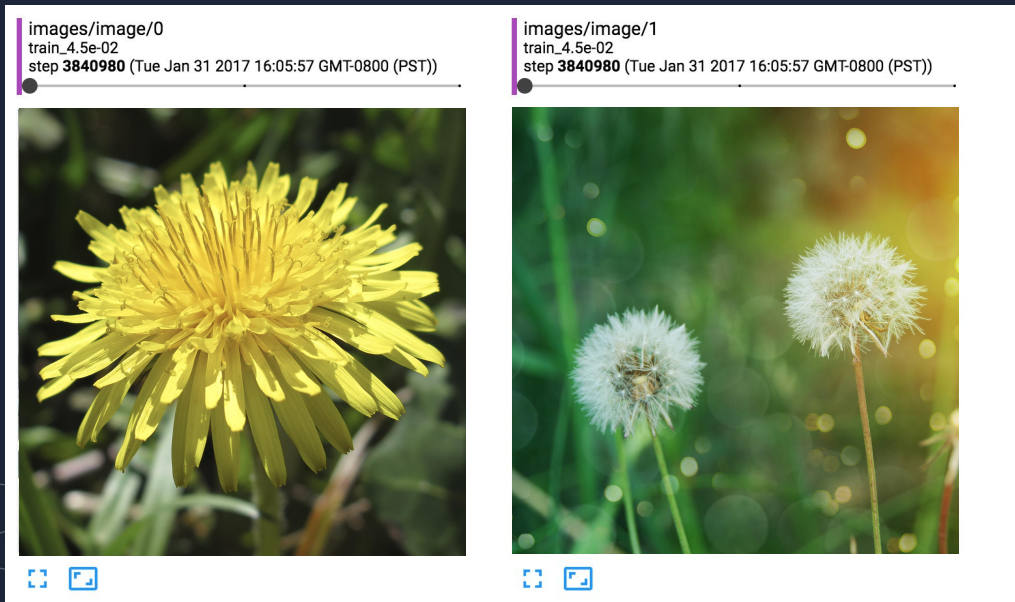
Collect some summaries

`summary (n):`

a TensorFlow op that output protocol buffers containing "summarized" data

Examples:

- `tf.summary.scalar`
- `tf.summary.image`



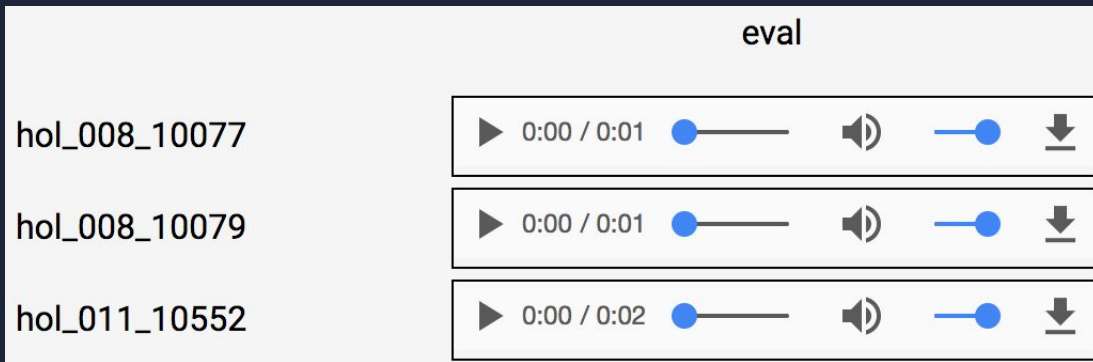
Collect some summaries

`summary (n):`

a TensorFlow op that output protocol buffers containing "summarized" data

Examples:

- `tf.summary.scalar`
- `tf.summary.image`
- **`tf.summary.audio`**



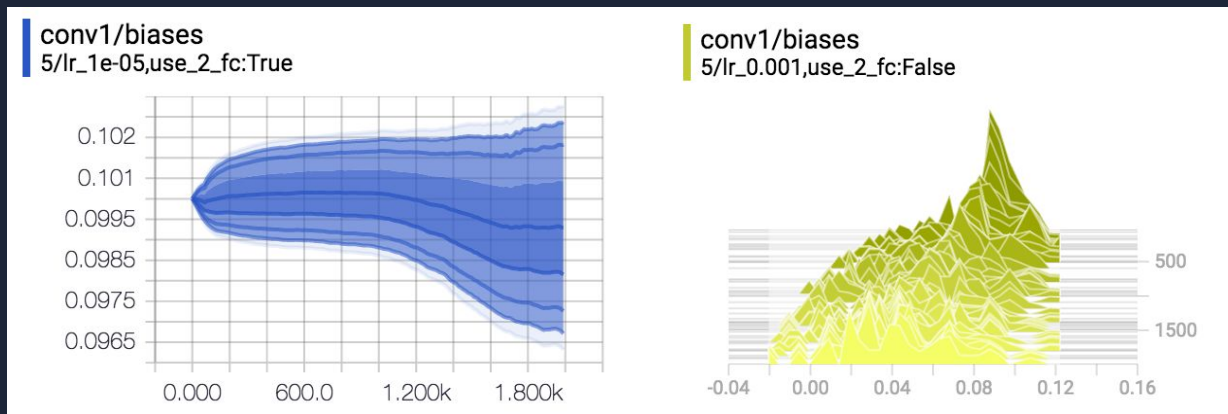
Collect some summaries

`summary (n):`

a TensorFlow op that output protocol buffers containing "summarized" data

Examples:

- `tf.summary.scalar`
- `tf.summary.image`
- `tf.summary.audio`
- **`tf.summary.histogram`**



Collect some summaries

`summary (n):`

a TensorFlow op that output protocol buffers containing "summarized" data

Examples:

- `tf.summary.scalar`
- `tf.summary.image`
- `tf.summary.audio`
- `tf.summary.histogram`
- **`tf.summary.tensor`**



Collect some summaries

`summary (n):`

a TensorFlow op that output protocol buffers containing "summarized" data

Examples:

- `tf.summary.scalar`
- `tf.summary.image`
- `tf.summary.audio`
- `tf.summary.histogram`
- **`tf.summary.tensor` (under development)**



Collect some summaries

```
tf.summary.scalar('cross_entropy', xent)  
tf.summary.scalar('accuracy', accuracy)
```



Collect some summaries

```
tf.summary.scalar('cross_entropy', xent)  
tf.summary.scalar('accuracy', accuracy)
```

```
tf.summary.image('input', x_image, 3)
```



Collect some summaries

```
def conv_layer(input, channels_in, channels_out, name="conv"):
    with tf.name_scope(name):
        w = tf.Variable(tf.zeros([5, 5, channels_in, channels_out]), name="W")
        b = tf.Variable(tf.zeros([channels_out]), name="B")
        conv = tf.nn.conv2d(input, w, strides=[1, 1, 1, 1], padding="SAME")
        act = tf.nn.relu(conv + b)
        tf.summary.histogram("weights", w)
        tf.summary.histogram("biases", b)
        tf.summary.histogram("activations", act)
    return tf.nn.max_pool(act, ksize=[1, 2, 2, 1], strides=[1, 2, 2, 1], padding="SAME")
```



How Summaries Work

- Summary Op returns protocol buffers
- `tf.summary.FileWriter` writes them to disk



Collect some summaries

```
merged_summary = tf.summary.merge_all()
```



Collect some summaries

```
merged_summary = tf.summary.merge_all()  
writer = tf.summary.FileWriter("/tmp/mnist_demo/3")  
writer.add_graph(sess.graph)
```



Collect some summaries

```
merged_summary = tf.summary.merge_all()
writer = tf.summary.FileWriter("/tmp/mnist_demo/3")
writer.add_graph(sess.graph)

for i in range(2001):
    batch = mnist.train.next_batch(100)
    if i % 5 == 0:
        s = sess.run(merged_summary, feed_dict={x: batch[0], y: batch[1]})
        writer.add_summary(s, i)
    sess.run(train_step, feed_dict={x: batch[0], y: batch[1]})
```



Collect some summaries

```
merged_summary = tf.summary.merge_all()
writer = tf.summary.FileWriter("/tmp/mnist_demo/3")
writer.add_graph(sess.graph)

for i in range(2001):
    batch = mnist.train.next_batch(100)
    if i % 5 == 0:
        s = sess.run(merged_summary, feed_dict={x: batch[0], y: batch[1]})
        writer.add_summary(s, i)
    sess.run(train_step, feed_dict={x: batch[0], y: batch[1]})
```

>> tensorboard --logdir /tmp/mnist_demo/3



(DEMO)



Fixing Our Model

```
def conv_layer(input, size_in, size_out, name="conv"):
    with tf.name_scope(name):
        w = tf.Variable(tf.zeros([5, 5, size_in, size_out]), name="W")
        b = tf.Variable(tf.zeros([size_out]), name="B")
        conv = tf.nn.conv2d(input, w, strides=[1, 1, 1, 1], padding="SAME")
        act = tf.nn.relu(conv + b)
        tf.summary.histogram("weights", w)
        tf.summary.histogram("biases", b)
        tf.summary.histogram("activations", act)
        return tf.nn.max_pool(act, ksize=[1, 2, 2, 1], strides=[1, 2, 2, 1], padding="SAME")
```



Fixing Our Model

```
def conv_layer(input, size_in, size_out, name="conv"):
    with tf.name_scope(name):
        w = tf.Variable(tf.truncated_normal([5, 5, size_in, size_out], stddev=0.1), name="W")
        b = tf.Variable(tf.constant(0.1, shape=[size_out]), name="B")
        conv = tf.nn.conv2d(input, w, strides=[1, 1, 1, 1], padding="SAME")
        act = tf.nn.relu(conv + b)
        tf.summary.histogram("weights", w)
        tf.summary.histogram("biases", b)
        tf.summary.histogram("activations", act)
        return tf.nn.max_pool(act, ksize=[1, 2, 2, 1], strides=[1, 2, 2, 1], padding="SAME")
```



(DEMO)



Hyperparameter Search

- What about different learning rates?
- What about different model architectures?



Hyperparameter Search

```
# Try a few learning rates
for learning_rate in [1E-3, 1E-4, 1E-5]:

    # Try a model with fewer layers
    for use_two_fc in [True, False]:
        for use_two_conv in [True, False]:

            # Construct a hyperparameter string for each one (example: "lr_1E-3,fc=2,conv=2")
            hparam_str = make_hparam_string(learning_rate, use_two_fc, use_two_conv)

            writer = tf.summary.FileWriter("/tmp/mnist_tutorial/" + hparam_str)

            # Actually run with the new settings
            mnist(learning_rate, use_two_fully_connected_layers, use_two_conv_layers, writer)
```

>> tensorboard --logdir /tmp/mnist_tutorial



(DEMO)



Embedding Visualization

```
embedding = tf.Variable(tf.zeros([10000, embedding_size]), name="test_embedding")
assignment = embedding.assign(embedding_input)

config = tf.contrib.tensorboard.plugins.projector.ProjectorConfig()
embedding_config = config.embeddings.add()
embedding_config.tensor_name = embedding.name
embedding_config.sprite.image_path = os.path.join(LOG_DIR, 'sprite.png')
# Specify the width and height of a single thumbnail.
embedding_config.sprite.single_image_dim.extend([28, 28])
tf.contrib.tensorboard.plugins.projector.visualize_embeddings(writer, config)
```



Embedding Visualization

```
for i in range(2001):  
    batch = mnist.train.next_batch(100)  
    if i % 5 == 0:  
        [train_accuracy, s] = sess.run([accuracy, summ], feed_dict={x: batch[0], y: batch[1]})  
        writer.add_summary(s, i)  
    if i % 500 == 0:  
        sess.run(assignment, feed_dict={x: mnist.test.images, y_true: mnist.test.labels})  
        saver.save(sess, os.path.join(LOG_DIR, "model.ckpt"), i)  
    sess.run(train_step, feed_dict={x: batch[0], y_true: batch[1]})
```

```
>> tensorboard --logdir /tmp/mnist_tutorial
```



(DEMO)



Future for TensorBoard



Future for TensorBoard

- TensorFlow Debugger Integration
- Plugins



Validation example 1 ▾ colored by activations of cell 55 ▾

some freak of e


Activation value




Step 0

①

Cell 55

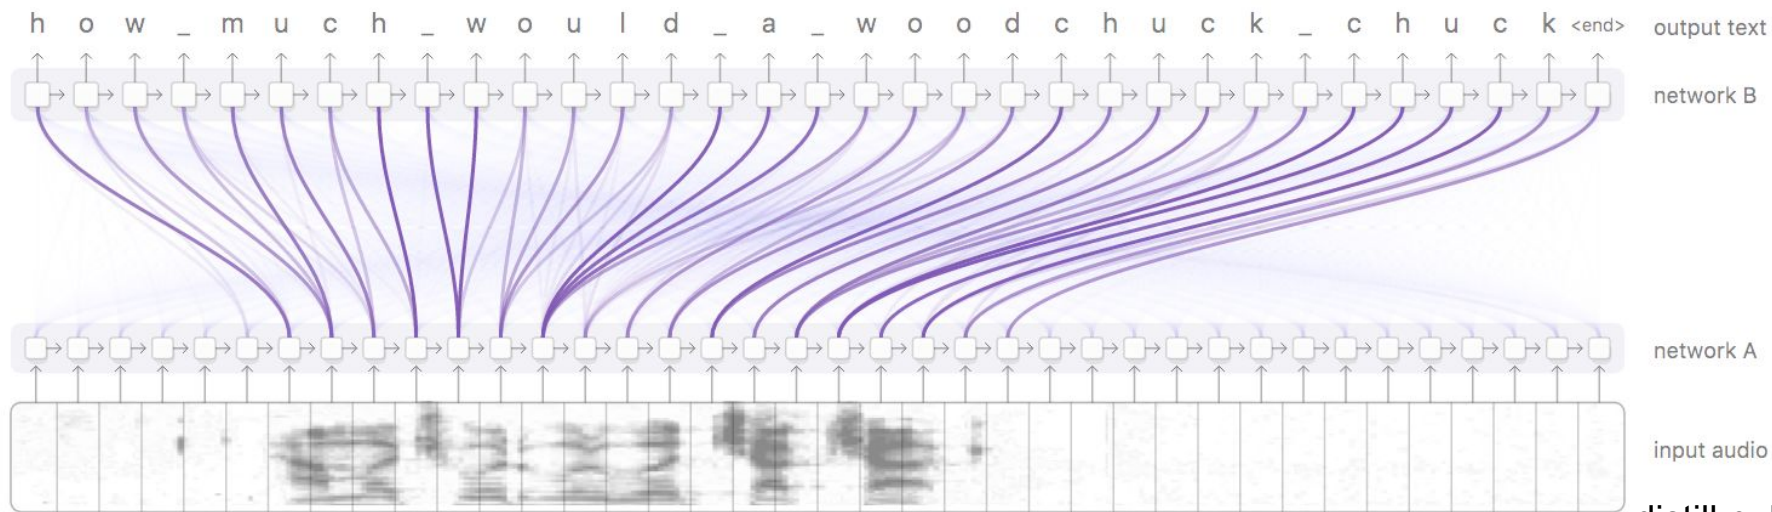


Activations at step 0



distill.pub





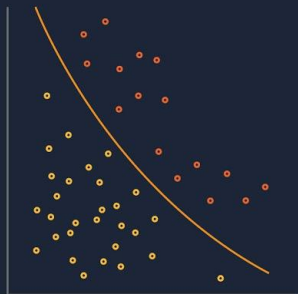
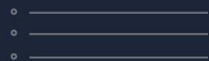
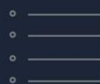
distill.pub



Future for TensorBoard

- TensorFlow Debugger Integration
- Plugins
- Org-scale TensorBoard





Code: <https://goo.gl/San2uR>



Slides: <https://goo.gl/4lfwZy>

